



CS 540 Introduction to Artificial Intelligence

Linear Algebra & PCA

Fred Sala
University of Wisconsin-Madison

Sep 15, 2022

Announcements

- **Homeworks:**
 - HW1 due 5 minutes ago; HW2 released today.
- **Class roadmap:**

Tuesday, Sep 13	Probability
Thursday, Sep 15	Linear Algebra and PCA
Tuesday, Sep 20	Statistics and Math Review
Thursday, Sep 22	Introduction to Logic
Tuesday, Sep 27	Natural Language Processing

} Fundamentals

From Last Time

- Conditional Prob. & Bayes:

$$P(H|E_1, E_2, \dots, E_n) = \frac{P(E_1, \dots, E_n|H)P(H)}{P(E_1, E_2, \dots, E_n)}$$

- Has more evidence.
 - Likelihood is hard---but **conditional independence assumption**

$$P(H|E_1, E_2, \dots, E_n) = \frac{P(E_1|H)P(E_2|H) \cdots P(E_n|H)P(H)}{P(E_1, E_2, \dots, E_n)}$$

Classification

- Expression

$$P(H|E_1, E_2, \dots, E_n) = \frac{P(E_1|H)P(E_2|H) \cdots P(E_n|H)P(H)}{P(E_1, E_2, \dots, E_n)}$$

- H : some class we'd like to infer from evidence
 - We know prior $P(H)$
 - Estimate $P(E_i|H)$ from data! (“training”)
 - Very similar to envelopes problem. **Part of HW2**

Linear Algebra: What is it good for?

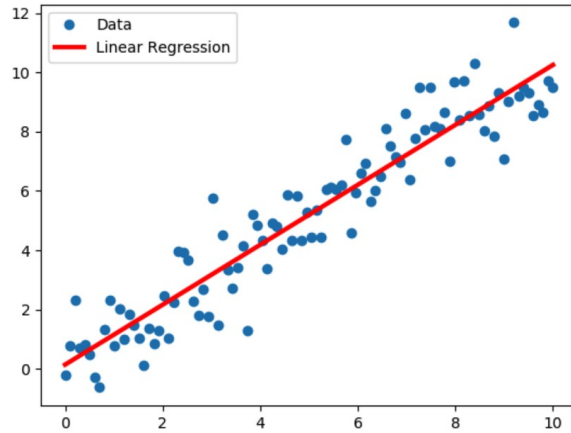
- Everything is a **function**
 - Multiple inputs and outputs
- Linear functions
 - Simple, tractable
- Study of linear functions



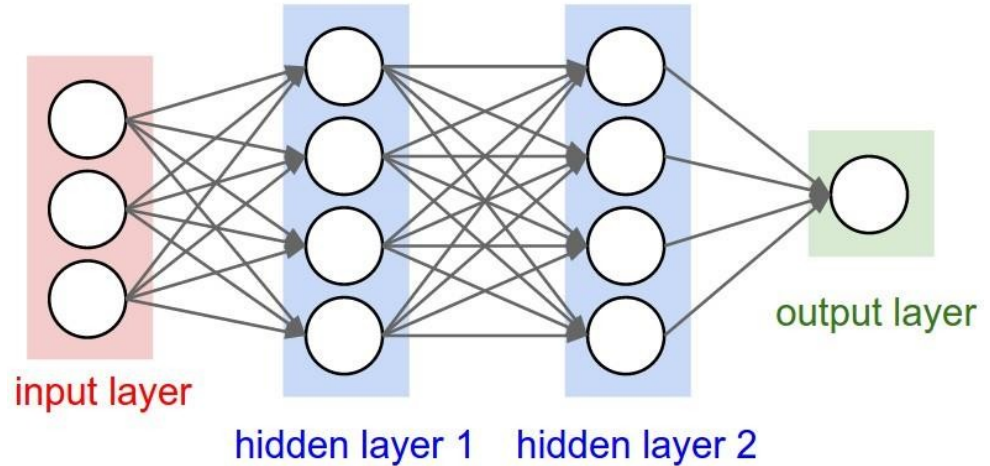
In AI/ML Context

Building blocks for **all** models

- E.g., linear regression; part of neural networks



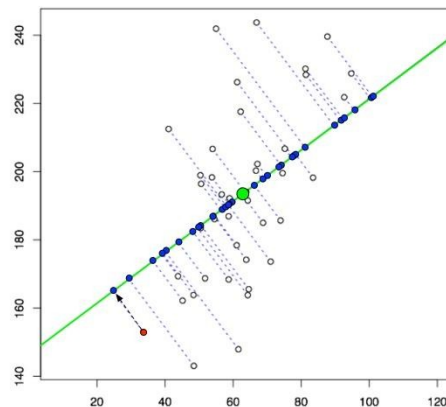
Hieu Tran



Stanford CS231n

Outline

- Basics: vectors, matrices, operations
- Dimensionality reduction
- Principal Components Analysis (PCA)

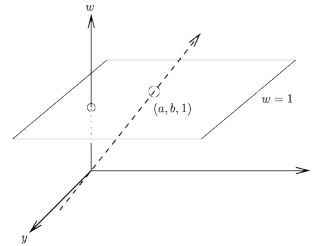
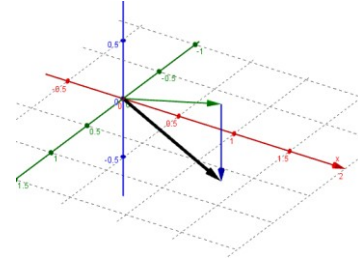


Lior Pachter

Basics: Vectors

Vectors

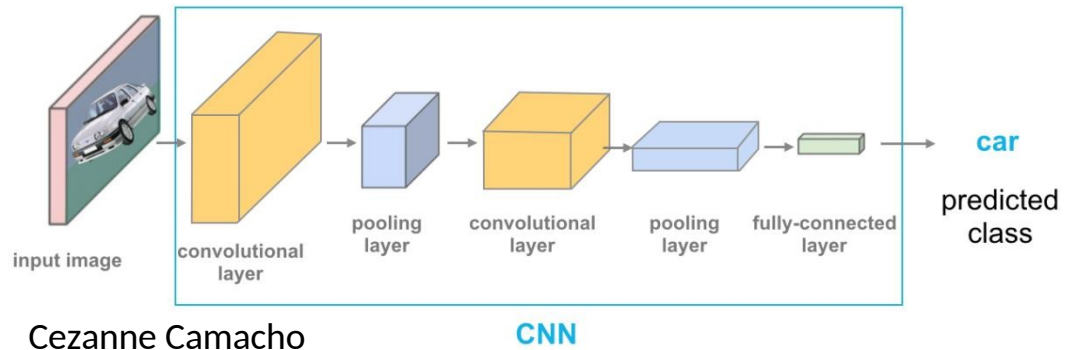
- Many interpretations
 - Physics: magnitude + direction
 - Point in a space
 - List of values (represents information)



$$x = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \end{bmatrix}$$

Basics: Vectors

- Dimension
 - Number of values $x \in \mathbb{R}^d$
 - Higher dimensions: richer but more complex
- AI/ML: often use **very high dimensions**:
 - Ex: images!



Basics: Matrices

- Again, many interpretations
 - Represent linear transformations
 - Apply to a vector, get another vector
 - Also, list of vectors

$$A = \begin{bmatrix} A_{11} & A_{12} & A_{13} \\ A_{21} & A_{22} & A_{23} \\ A_{31} & A_{32} & A_{33} \end{bmatrix}$$

- Not necessarily square
 - Indexing! $A \in \mathbb{R}^{c \times d}$
 - Dimensions: #rows x #columns

Basics: Transposition

- Transposes: flip rows and columns
 - Vector: standard is a column. Transpose: row
 - Matrix: go from $m \times n$ to $n \times m$

$$x = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} \quad x^T = \begin{bmatrix} x_1 & x_2 & x_3 \end{bmatrix}$$

$$A = \begin{bmatrix} A_{11} & A_{12} & A_{13} \\ A_{21} & A_{22} & A_{23} \end{bmatrix} \quad A^T = \begin{bmatrix} A_{11} & A_{21} \\ A_{12} & A_{22} \\ A_{13} & A_{23} \end{bmatrix}$$

Matrix & Vector Operations

- Vectors

- Addition: component-wise

- Commutative
 - Associative

$$x + y = \begin{bmatrix} x_1 + y_1 \\ x_2 + y_2 \\ x_3 + y_3 \end{bmatrix}$$

- Scalar Multiplication

- Uniform stretch / scaling

$$cx = \begin{bmatrix} cx_1 \\ cx_2 \\ cx_3 \end{bmatrix}$$

Matrix & Vector Operations

- **Vector products.**

- Inner product (e.g., dot product)

$$\langle x, y \rangle := x^T y = \begin{bmatrix} x_1 & x_2 & x_3 \end{bmatrix} \begin{bmatrix} y_1 \\ y_2 \\ y_3 \end{bmatrix} = x_1 y_1 + x_2 y_2 + x_3 y_3$$

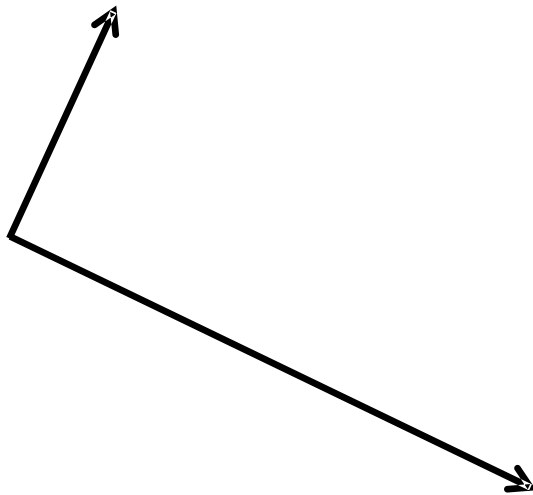
- Outer product

$$xy^T = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} \begin{bmatrix} y_1 & y_2 & y_3 \end{bmatrix} = \begin{bmatrix} x_1 y_1 & x_1 y_2 & x_1 y_3 \\ x_2 y_1 & x_2 y_2 & x_2 y_3 \\ x_3 y_1 & x_3 y_2 & x_3 y_3 \end{bmatrix}$$

Matrix & Vector Operations

- Inner product defines “orthogonality”
 - If $\langle x, y \rangle = 0$
- Vector norms: “size”

$$\|x\|_2 = \sqrt{\sum_{i=1}^n x_i^2}$$



Matrix & Vector Operations

- Matrices:

- Addition: Component-wise

- **Commutative!** + Associative

$$A + B = \begin{bmatrix} A_{11} + B_{11} & A_{12} + B_{12} \\ A_{21} + B_{21} & A_{22} + B_{22} \\ A_{31} + B_{31} & A_{32} + B_{32} \end{bmatrix}$$

- Scalar Multiplication

- “Stretching” the linear transformation

$$cA = \begin{bmatrix} cA_{11} & cA_{12} \\ cA_{21} & cA_{22} \\ cA_{31} & cA_{32} \end{bmatrix}$$

Matrix & Vector Operations

- Matrix-Vector multiply
 - I.e., linear transformation; plug in vector, get another vector
 - Each entry in Ax is the inner product of a row of A with x

$$Ax = \begin{bmatrix} A_{11}x_1 + A_{12}x_2 + \dots + A_{1n}x_n \\ A_{21}x_1 + A_{22}x_2 + \dots + A_{2n}x_n \\ \vdots \\ A_{n1}x_1 + A_{n2}x_2 + \dots + A_{nn}x_n \end{bmatrix}$$

Matrix & Vector Operations

Ex: feedforward neural networks. Input x .

- Output of layer k is

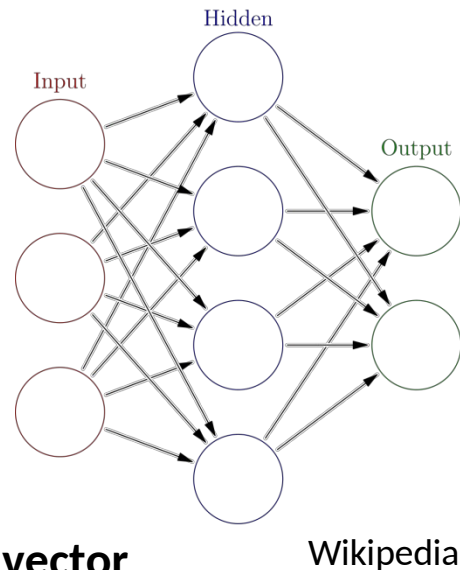
$$f^{(k)}(x) = \sigma(W_k^T f^{(k-1)}(x))$$

nonlinearity

Output of layer $k-1$: **vector**

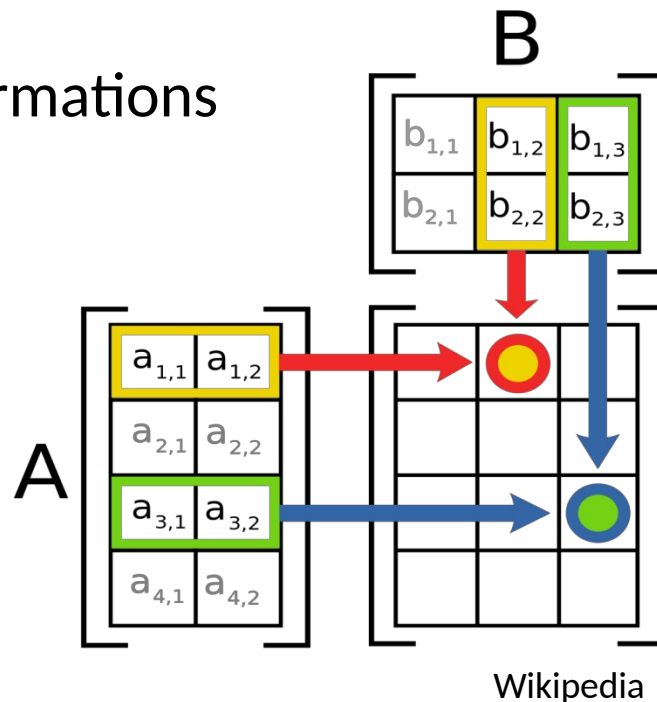
Output of layer k : vector

Weight **matrix** for layer k :
Note: linear transformation!



Matrix & Vector Operations

- Matrix multiplication
 - “Composition” of linear transformations
 - **Not commutative** (in general)!
 - Lots of interpretations



More on Matrix Operations

- Identity matrix:
 - Like “1”
 - Multiplying by it gets back the same matrix or vector
 - Rows & columns are the “**standard basis vectors**” e_i

$$I = \begin{bmatrix} 1 & 0 & \dots & 0 \\ 0 & 1 & \dots & 0 \\ \vdots & \vdots & \dots & \vdots \\ 0 & 0 & \dots & 1 \end{bmatrix}$$

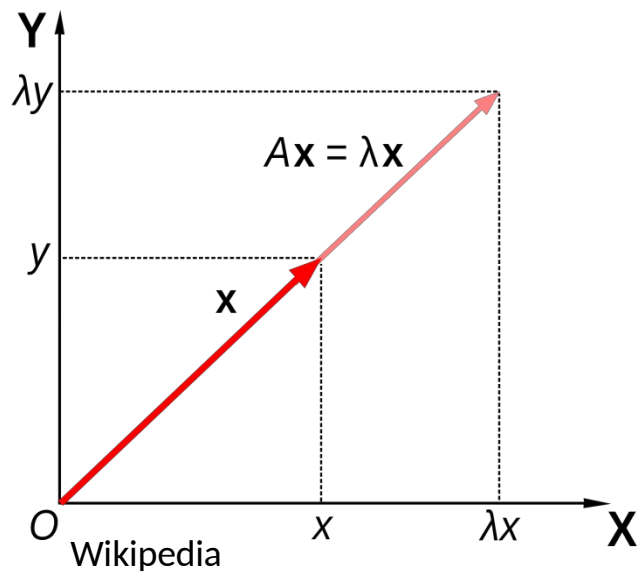
More on Matrices: Inverses

- If for A there is a B such that $AB = BA = I$
 - Then A is invertible/nonsingular, B is its inverse
 - Some matrices are **not** invertible!
 - Usual notation: A^{-1}

$$\begin{bmatrix} 1 & 1 \\ 2 & 3 \end{bmatrix} \times \begin{bmatrix} 3 & -1 \\ -2 & 1 \end{bmatrix} = I$$

Eigenvalues & Eigenvectors

- For a square matrix A , solutions to $Av = \lambda v$
 - v (nonzero) is a vector: **eigenvector**
 - λ is a scalar: **eigenvalue**
 - Intuition: A is a linear transformation;
 - Can stretch/rotate vectors;
 - E-vectors: only stretched (by e-vals)



Dimensionality Reduction

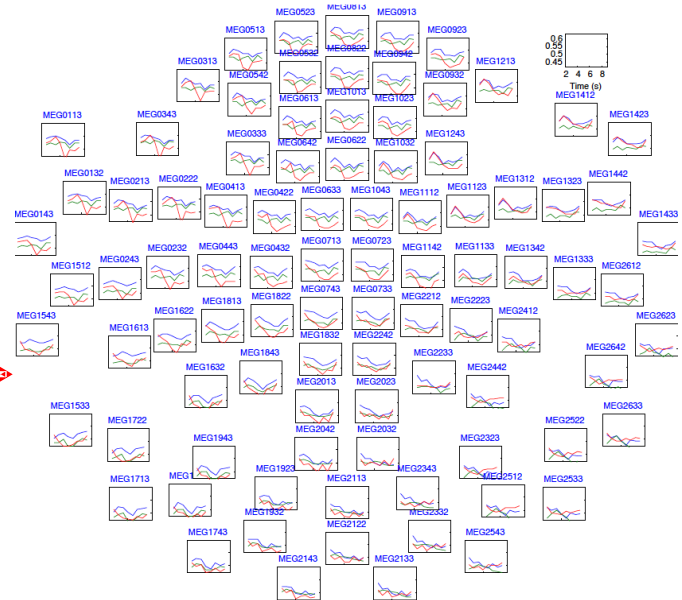
- Vectors used to store features
 - Lots of data -> lots of features!
- Document classification
 - Each doc: thousands of words/millions of bigrams, etc
- Netflix surveys: 480189 users x 17770 movies

	movie 1	movie 2	movie 3	movie 4	movie 5	movie 6
Tom	5	?	?	1	3	?
George	?	?	3	1	2	5
Susan	4	3	1	?	5	1
Beth	4	3	?	2	4	2

Dimensionality Reduction

Ex: **MEG Brain Imaging**: 120 locations x 500 time points
x 20 objects

- Or any image



Dimensionality Reduction

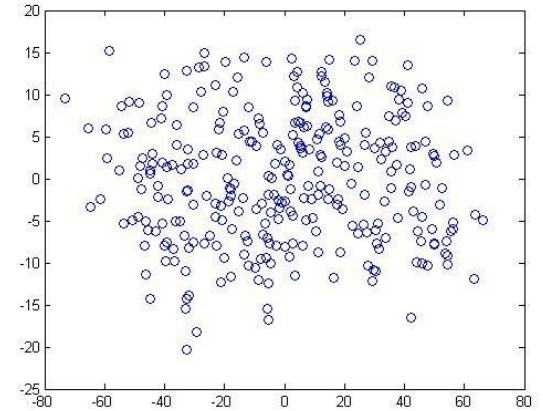
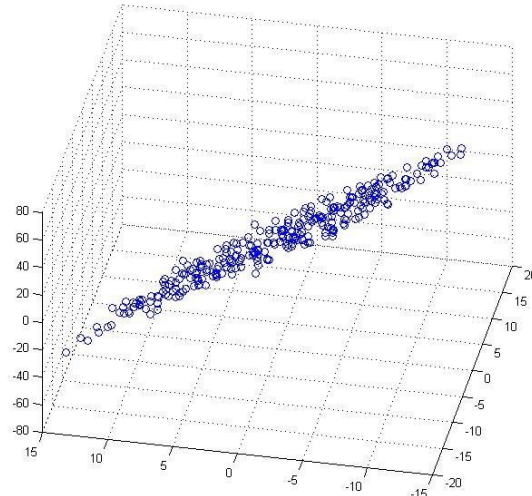
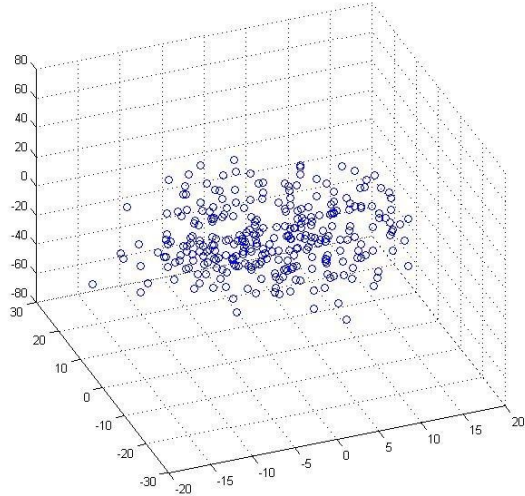
Reduce dimensions

- Why?
 - Lots of features redundant
 - Storage & computation costs
- Goal: take $x \in \mathbb{R}^d \rightarrow x \in \mathbb{R}^r$ for $r \ll d$
 - But, minimize information loss



Compression

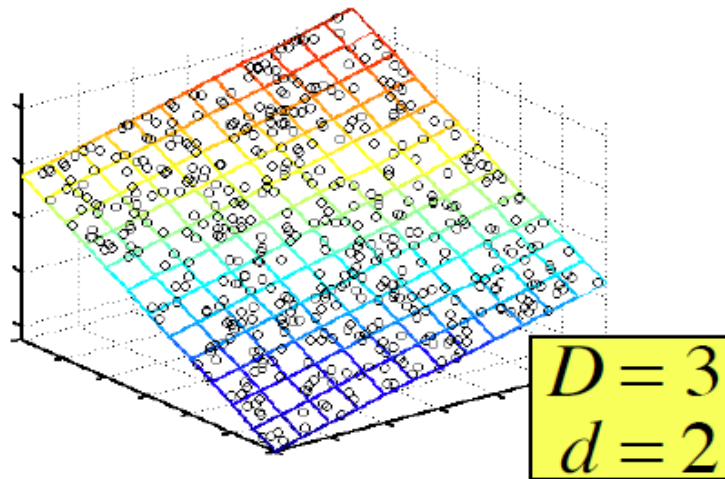
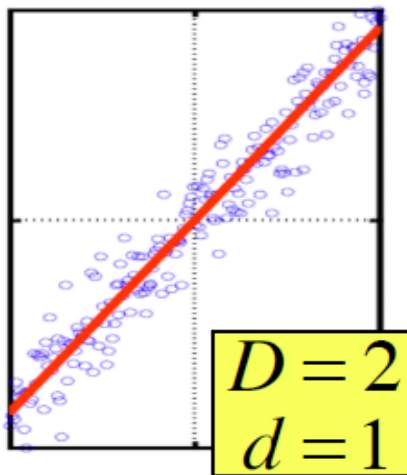
Examples: 3D to 2D



Andrew Ng

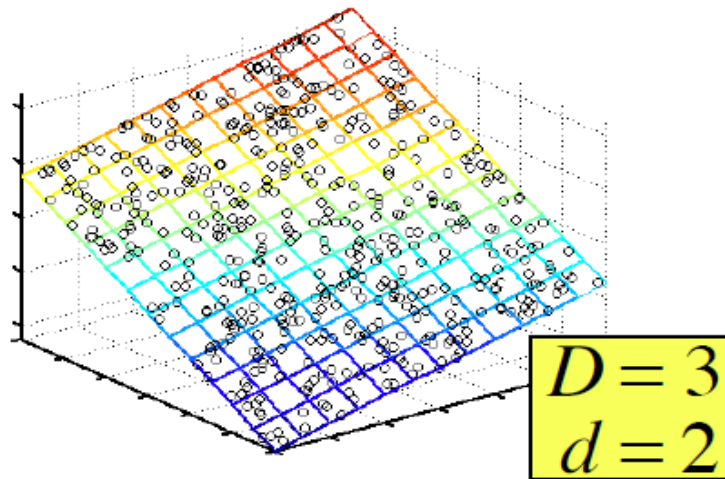
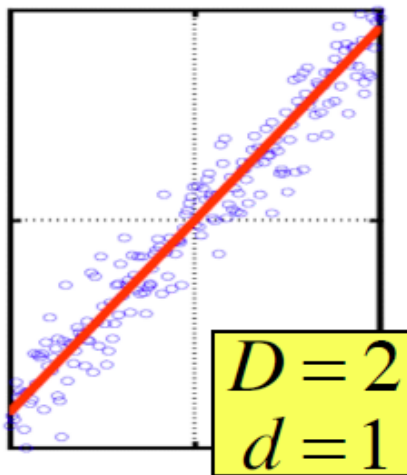
Principal Components Analysis (PCA)

- A type of dimensionality reduction approach
 - For when data is **approximately lower dimensional**



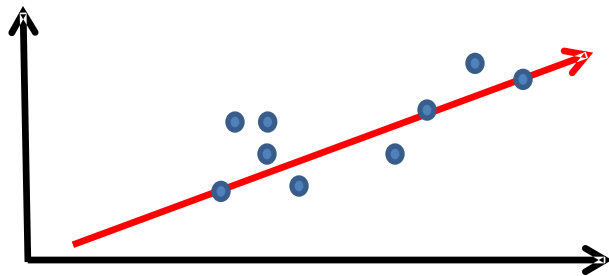
Principal Components Analysis (PCA)

- Goal: find **axes** of a subspace
 - Will project to this subspace; want to preserve data



Principal Components Analysis (PCA)

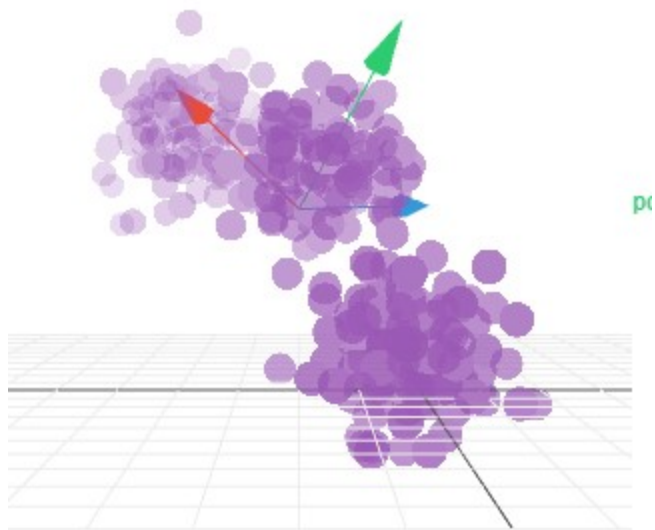
- From 2D to 1D:
 - Find a $v_1 \in \mathbb{R}^d$ so that we maximize “variability”
 - IE,



- New representations are along this vector (1D!)

Principal Components Analysis (PCA)

- From d dimensions to r dimensions
 - Sequentially get $v_1, v_2, \dots, v_r \in \mathbb{R}^d$
 - Orthogonal!
 - Still minimize the projection error
 - Equivalent to “**maximizing variability**”
 - The vectors are the **principal components**



Victor Powell

PCA Setup

- **Inputs**

- Data: $x_1, x_2, \dots, x_n, x_i \in \mathbb{R}^d$

- Can arrange into $X \in \mathbb{R}^{n \times d}$

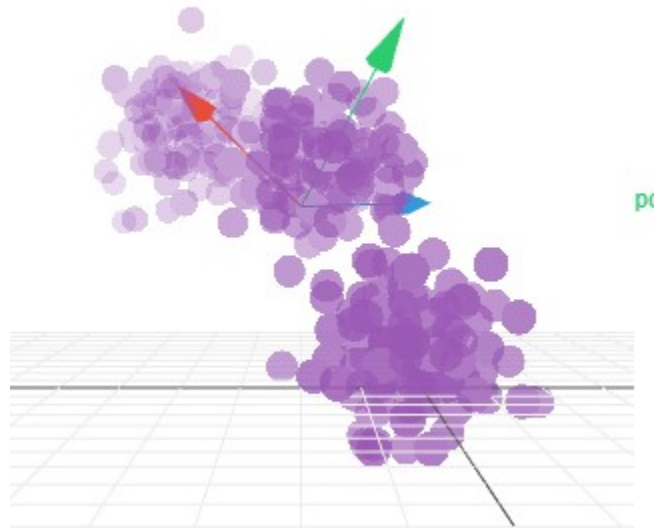
- **Centered!**

$$\frac{1}{n} \sum_{i=1}^n x_i = 0$$

- **Outputs**

- Principal components $v_1, v_2, \dots, v_r \in \mathbb{R}^d$

- Orthogonal!



Victor Powell

PCA Goals

- Want directions/components (unit vectors) so that
 - Projecting data maximizes variance
 - What's projection?

$$\sum_{i=1}^n \langle x_i, v \rangle = \|Xv\|^2$$

- Do this **recursively**
 - Get orthogonal directions $v_1, v_2, \dots, v_r \in \mathbb{R}^d$

PCA First Step

- First component,

$$v_1 = \arg \max_{\|v\|=1} \sum_{i=1}^n \langle v, x_i \rangle^2$$

- Same as getting

$$v_1 = \arg \max_{\|v\|=1} \|Xv\|^2$$

PCA Recursion

- Once we have $k-1$ components, next?

$$\hat{X}_k = X - \sum_{i=1}^{k-1} X v_i v_i^T$$

- Then do the same thing

Deflation



$$v_k = \arg \max_{\|v\|=1} \|\hat{X}_k w\|^2$$

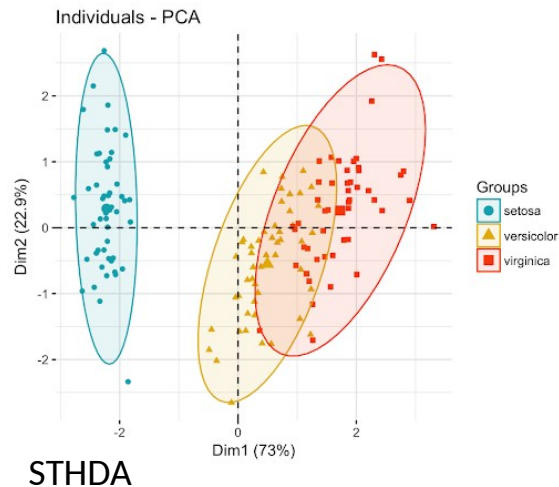
PCA Interpretations

- The v 's are eigenvectors of $X^T X$ (**Gram matrix**)
 - Show via Rayleigh quotient
- $X^T X$ (proportional to) sample covariance matrix
 - When data is 0 mean!
 - I.e., PCA is eigendecomposition of sample covariance
- Nested subspaces $\text{span}(v_1)$, $\text{span}(v_1, v_2)$, ...,



Lots of Variations

- PCA, Kernel PCA, ICA, CCA
 - Unsupervised techniques to extract structure from high dimensional dataset
- Uses:
 - **Visualization**
 - Efficiency
 - Noise removal
 - Downstream machine learning use



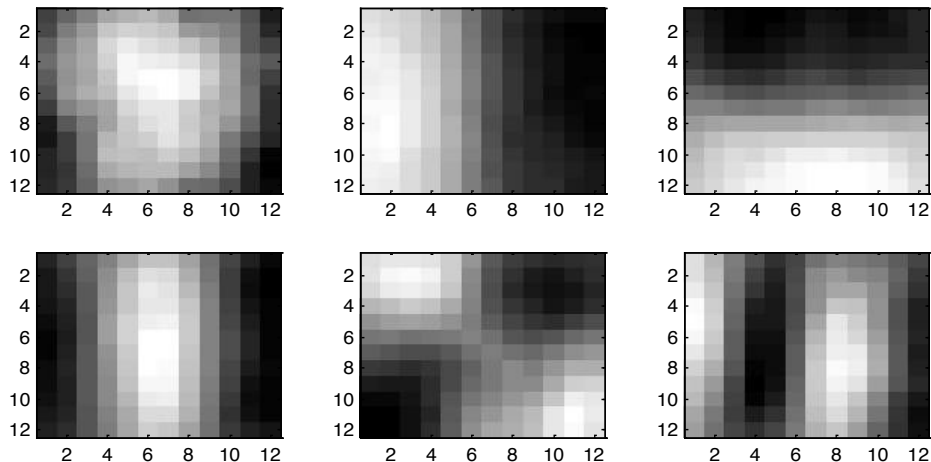
Application: Image Compression

- Start with image; divide into 12x12 patches
 - I.E., 144-D vector
 - **Original image:**



Application: Image Compression

- 6 most important components (as an image)



Application: Image Compression

- Project to 6D,



Compressed



Original